



# Introduction to Docker and Containers

CS 4230

Jay Urbain, Ph.D.

**YOU GET A CLUSTER AND YOU GET A CLUSTER**



**EVERYBODY GETS A CLUSTER**

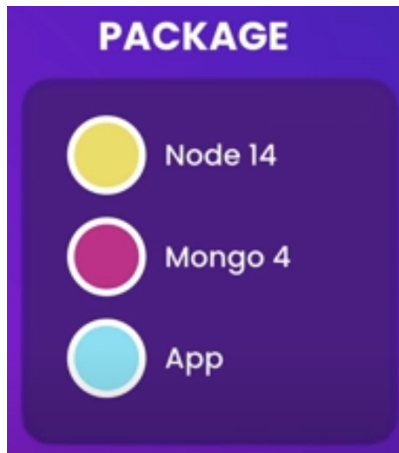
# Problems with application deployment

- One or more files missing
- Software version mismatch on target machine.
- Different configuration settings like environment variables.



# Docker

- A platform for building, running and deploying applications in a consistent manner.
- Create a package (container) of your application environment.
- If the package runs on your development machine, it should work on any target machine.



# Docker - additional benefits

- New developer don't have to spend a lot of time configuring a development environment.

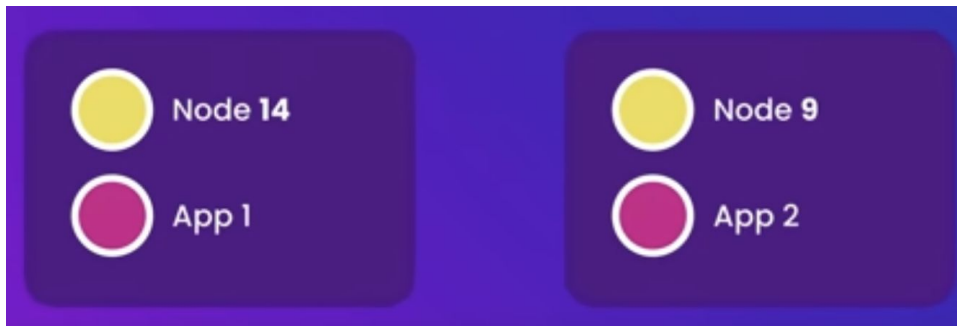
```
$ docker -compose up
```

Docker will download and run an application environment in an isolated environment called a **container**.



# Docker - additional benefits

- Isolated environment allows multiple applications to use different versions of the same software side by side.
- Also, when you're done with a specific version you can easily remove it in one shot.
  - You don't accumulate multiple versions of libraries, tools and environments on your development system.
  - You don't have to worry if it's safe to remove different development artifacts.
  - `$ docker -compose down -rm all`



# Containers versus Virtual Machines

Advantages/disadvantages of each approach?

## **CONTAINER**

An isolated environment for  
running an application

## **VIRTUAL MACHINE**

An abstraction of a machine  
(physical hardware)

# Virtual Machines

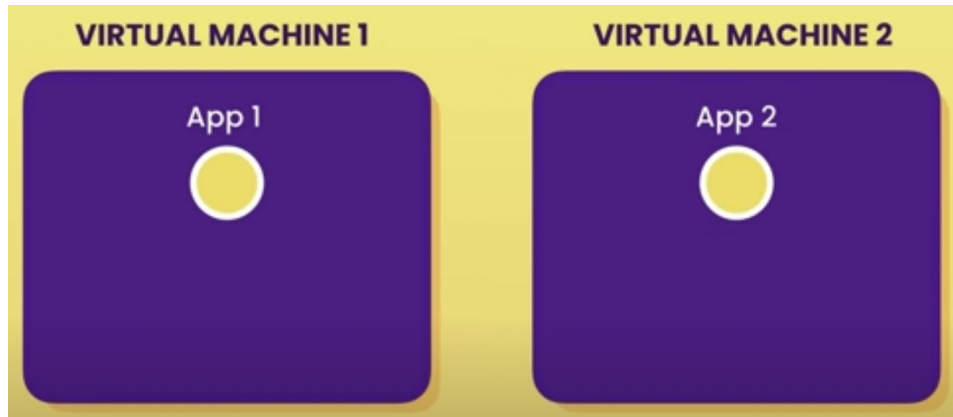
- Hypervisor - software to create and manage virtual machines.
- VirtualBox, VMware, Hyper-V (Windows)





# Virtual Machines

- Can run an application in near complete isolation inside a virtual machine on the same physical hardware.
- Needs a full copy of OS that may need to be licensed and patched.
- More resource intensive because each virtual machine takes a slice of physical hardware and must load the entire OS.



# Virtual Machines

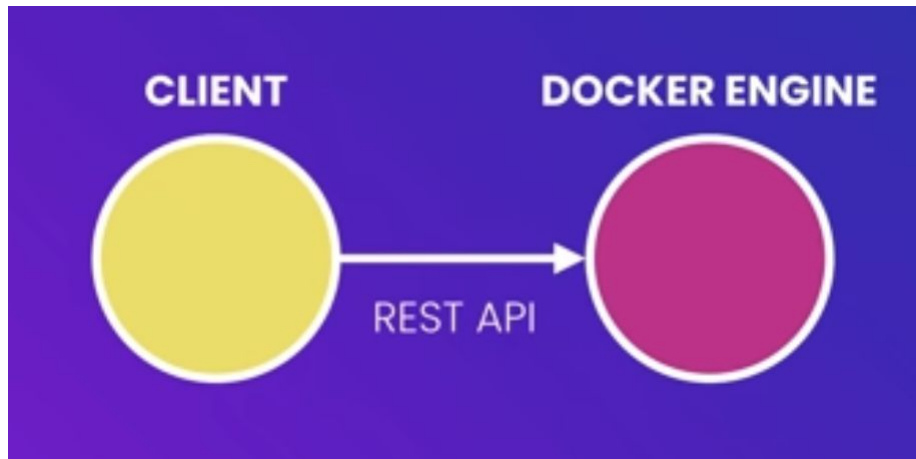
- Can run an application in isolation (exact dependencies) inside a virtual machine on the same physical hardware.
- Needs a full copy of OS that may need to be licensed and patched.
- Slow to start because entire OS had to be loaded.
- More resource intensive because each virtual machine takes a slice of physical hardware.
- Memory has to be divided among virtual machines.

# Containers

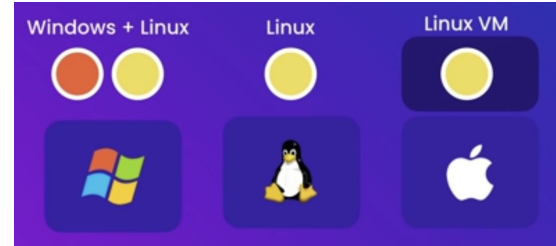
- Allow running multiple apps in isolation.
- More lightweight.
- Uses the OS of the host.
- Starts quickly - doesn't have to load OS.
- Need less hardware resources - don't need to allocate specific number of CPU cores or a slice of memory like VM.

# Docker

- Client-server architecture
- Docker engine runs in the background and takes care of building and running docker containers.
- Technically, a docker container is just a process like other processes running on your computer.



# Docker

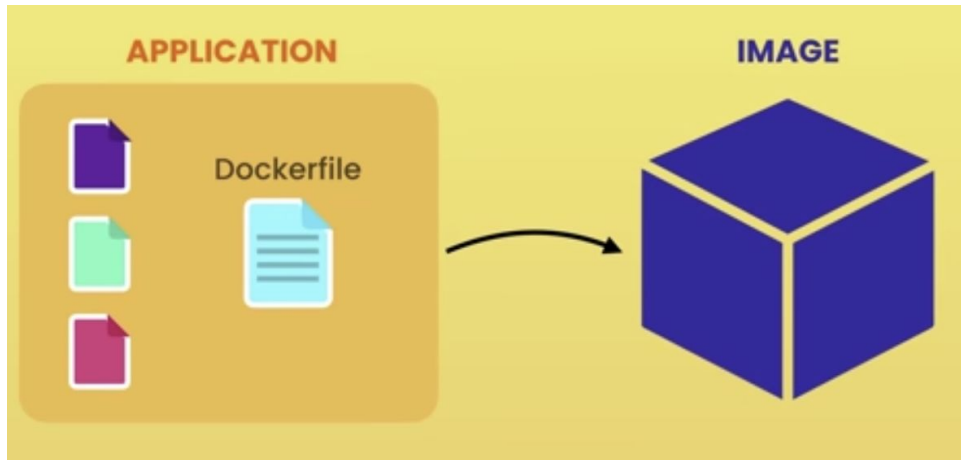


- Containers share the same operating system.
- Specifically the same OS kernel that manages hardware resources like memory, cpu, and scheduling.
- Different kernels have different APIs, that's why you can't run a Windows application on Linux.
- So on a Linux host, you can only run Linux containers.
- Can run both windows and Linux containers on Windows 10 since Windows now has a custom built linux kernel

# Docker Workflow

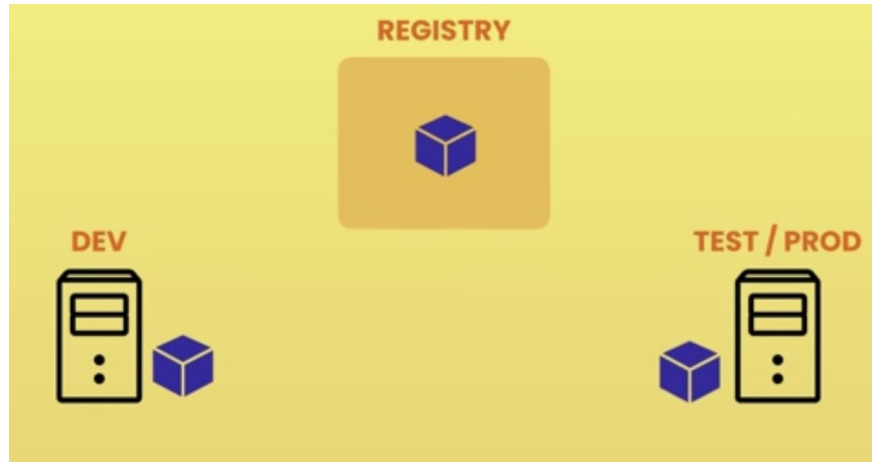
Dockerize your application.

- Dockerfile contains instructions for installing and running your application.
- Minified OS
- Runtime environment, e.g., Node or Python
- Application files
- Third party libraries/packages
- Environment variables

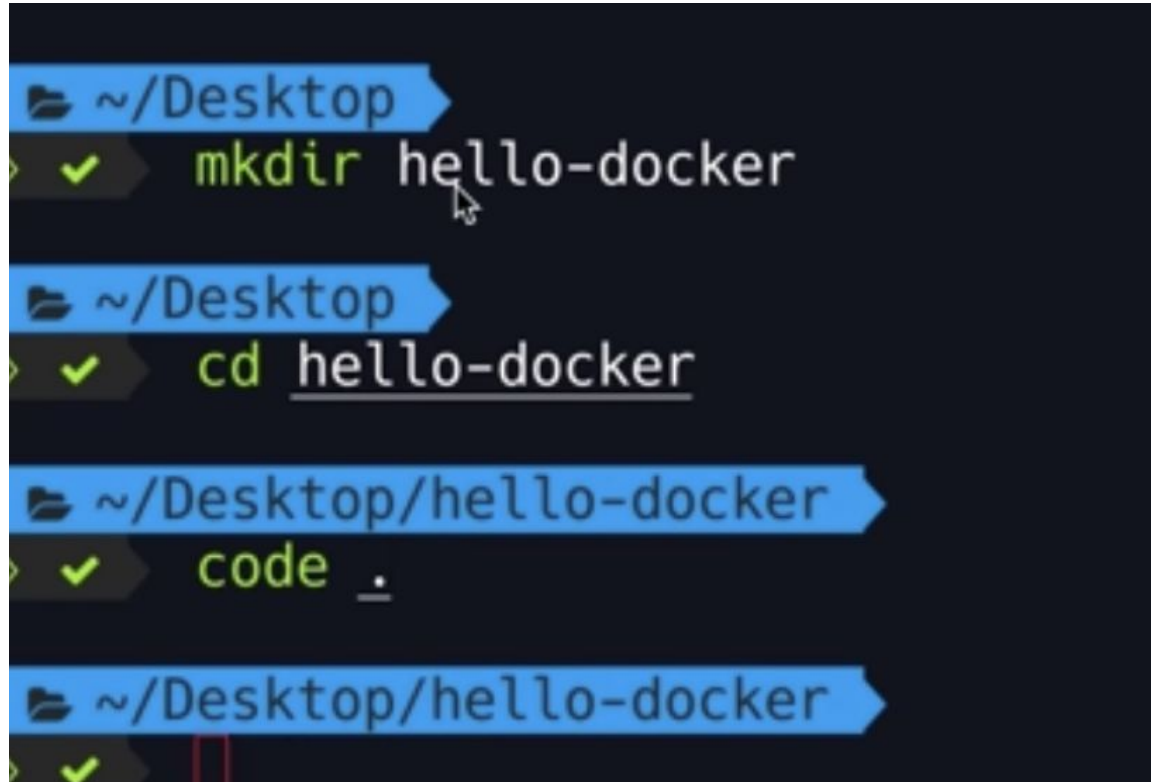


# Docker Workflow

- Special process since it contains its own file system which is provided by the image.
- Test/Prod has the same environment as our development machine.



## Docker Workflow - create an application directory



```
~/Desktop  
✓ mkdir hello-docker  
  
~/Desktop  
✓ cd hello-docker  
  
~/Desktop/hello-docker  
✓ code .  
  
~/Desktop/hello-docker  
✓
```

The image shows a terminal window with a dark background. It displays a sequence of four terminal sessions, each with a directory path in a blue header bar and a command in a green prompt. The first session is at ~/Desktop, where the command 'mkdir hello-docker' is entered. The second session is also at ~/Desktop, where the command 'cd hello-docker' is entered, with 'hello-docker' underlined. The third session is at ~/Desktop/hello-docker, where the command 'code .' is entered, with '.' underlined. The fourth session is also at ~/Desktop/hello-docker, showing a green prompt and a red cursor. Each command is preceded by a green checkmark icon.



# Docker Workflow - write an application

A screenshot of a code editor window titled "app.js — hello-docker". The editor has a dark theme. The top bar shows a file named "app.js" with a close button. The main editing area contains the code `console.log("Hello Docker!");` with syntax highlighting. On the left side, there is a vertical toolbar with icons for file explorer, search, source control, and other development tools.

# Docker Workflow - create a Docker image

## INSTRUCTIONS

- Start with an OS
- Install Node
- Copy app files
- Run node app.js

# Docker Workflow

Write instructions for creating and running image within Dockerfile

```
Dockerfile •  
Dockerfile > WORKDIR  
FROM node:alpine  
COPY . /app  
WORKDIR /app  
CMD node app.js
```

# Docker Workflow

## Build the image

```
~/Desktop/hello-docker
✓ docker build -t hello-docker .
+] Building 1.2s (8/8) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 99B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/node:alpine
=> [internal] load build context
=> => transferring context: 158B
=> CACHED [1/3] FROM docker.io/library/node:alpine@sha256:c01b5723
=> [2/3] COPY . /app
=> [3/3] WORKDIR /app
=> exporting to image
=> => exporting layers
```

# Docker Workflow

List and run image

```
~/Desktop/hello-docker
```

```
> ✓ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hello-docker	latest	8872d3105639	2 minutes ago	112MB

```
~/Desktop/hello-docker
```

```
> ✓ docker run hello-docker
```

```
Hello Docker!
```

# Docker Workflow

```
$docker pull jayurbain/deepnlpintro
```

# Docker Installation

Check your version of Docker. You want the latest version.

```
% docker --version
```


Docker version 20.10.23, build 7155243

## Download

<https://www.docker.com/get-started/>

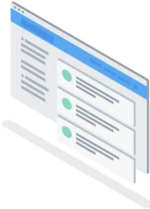
### Get Started with Docker

We have a complete container solution for you – no matter who you are and where you are on your containerization journey.



**Docker Desktop**  
Developer productivity tools and a local Kubernetes environment.

Download for Mac - Intel Chip



**Docker Hub**  
Cloud-based application registry and development team collaboration services.

Sign up

# Docker Installation

## **Docker Desktop**

Docker Desktop is an application for MacOS and Windows machines for the building and sharing of containerized applications.

## **Docker Hub**

Docker Hub is like github for containers.



# Docker Installation

Windows users: Make sure you enable Hyper-V on windows machine:

## System Requirements

- Windows 10 64-bit: Pro, Enterprise, or Education (Build 17134 or later).

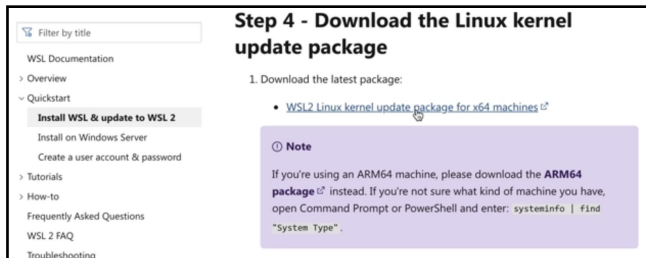
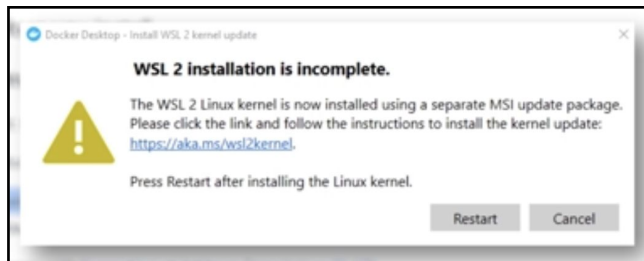
For Windows 10 Home, see [Install Docker Desktop on Windows Home](#).

- **Hyper-V and Containers Windows features** must be enabled.
- The following hardware prerequisites are required to successfully run Client Hyper-V on Windows 10:
  - 64 bit processor with [Second Level Address Translation \(SLAT\)](#)
  - 4GB system RAM
  - BIOS-level hardware virtualization support must be enabled in the BIOS settings. For more information, see [Virtualization](#).

# Docker Installation

Windows users:

If the following error occurs, you have to update your Windows Linux kernel that is shipped with your version of Windows.



# Run Docker

Make sure Docker is running. You should see the following:

```
% docker version
```

Client:

Cloud integration: v1.0.31

Version: 20.10.23

API version: 1.41

Go version: go1.18.10

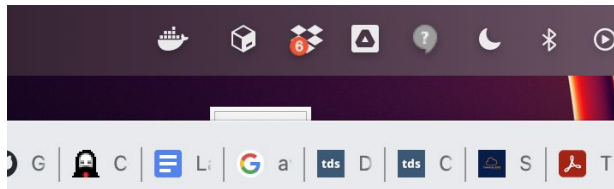
Git commit: 7155243

Built: Thu Jan 19 17:35:19 2023

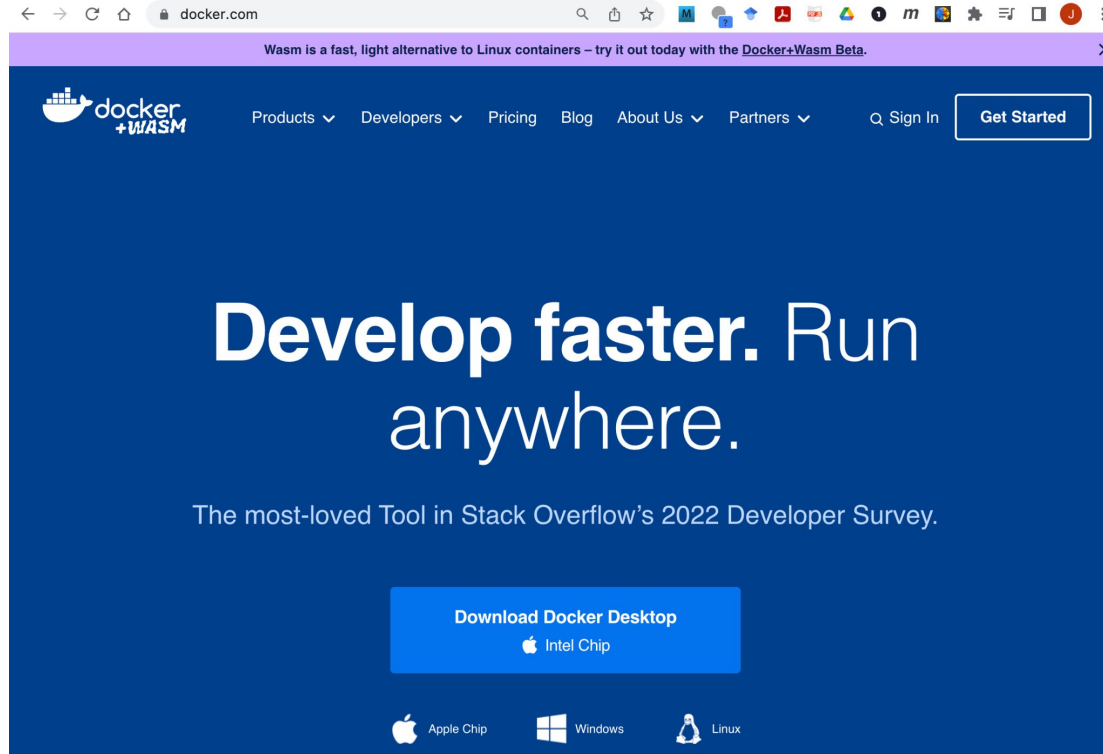
OS/Arch: darwin/arm64

...

You should also see the whale holding containers in your status bar. This shows the Docker Engine (server is running).



# Complete the Get Started Tutorial



# Docker pull

Bonus: Download and log into a Docker container from Docker Hub.