



Map Reduce

CS4230

Jay Urbain, Ph.D.

Credits:

- Jeffery Dean and Sanjay Chemawat. *MapReduce*
- Barroso and Urs Hölzle (2009)
- Jimmy Lin and Chris Dyer. *Data Intensive Text Processing with MapReduce*



The datacenter *is* the computer

Where in the datacenter are the bottlenecks in processing large data workloads?

Big Ideas

- **Scale *out*, not *up***
 - Limits of SMP (symmetric multi-processing) and large shared-memory machines
- **Move processing to the data**
 - Cluster has limited bandwidth
- **Process data sequentially, avoid random access**
 - Seeks are expensive, disk throughput is reasonable
 - Even with SSD sequential data is more likely to be co-located.
- **Seamless scalability**
 - Without synchronization issues, we move from the mythical man-month to the tradable machine-hour







Typical Large-Data Problem

- Iterate over a large number of records
- Extract something of interest from each
- Shuffle and sort intermediate results
- Aggregate intermediate results
- Generate final output

Map


Reduce

Key idea: provide a functional abstraction for these two operations.

Roots in Functional Programming

Functional programming

- Treats computation as the evaluation of mathematical functions.
- Avoids state and mutable data.
- Map and fold: map takes function f applies it to every element in a list, fold iteratively applies g to aggregate results.



MapReduce – Map + Reduce

1. "Map" step:

- The master node takes the input,
- partitions it up into smaller sub-problems,
- distributes to worker nodes.
- A worker node may do this again in turn, leading to a multi-level tree structure.
- The worker node processes that smaller problem, and passes the answer back to its master node.

2. "Reduce" step:

- The master node then takes the answers to all the sub-problems and *combines* them in some way to get the output — the answer to the problem it was originally trying to solve.

MapReduce - Programmatic

- Programmers specify two functions:

map $(k, v) \rightarrow \langle k', v' \rangle^*$


reduce $(k', v') \rightarrow \langle k', v' \rangle^*$

- Input data stored on underlying distributed file system.
- *Mapper* is applied to every input *key-value pair* (k, v) to generate an arbitrary number of intermediate *key-value pairs* $\langle k', v' \rangle^*$.
- *Reducer* is applied to all values associated with the same intermediate key (k', v') to generate output key-value pairs $\langle k', v' \rangle^*$.
- Implicit between the *map* and *reduce* phases is a distributed “*group by*” operation on intermediate keys.
- Intermediate data arrive at each reducer in key sorted order.


MapReduce - Notes

Notes:

- Output key-value pairs from each *reducer* are written persistently back onto the distributed file system.
- Intermediate key-values pairs are *transient* and are not preserved.
- The output ends up in r files on the distributed file system where r is the number of reducers.



Shuffle and Sort: aggregate values by keys



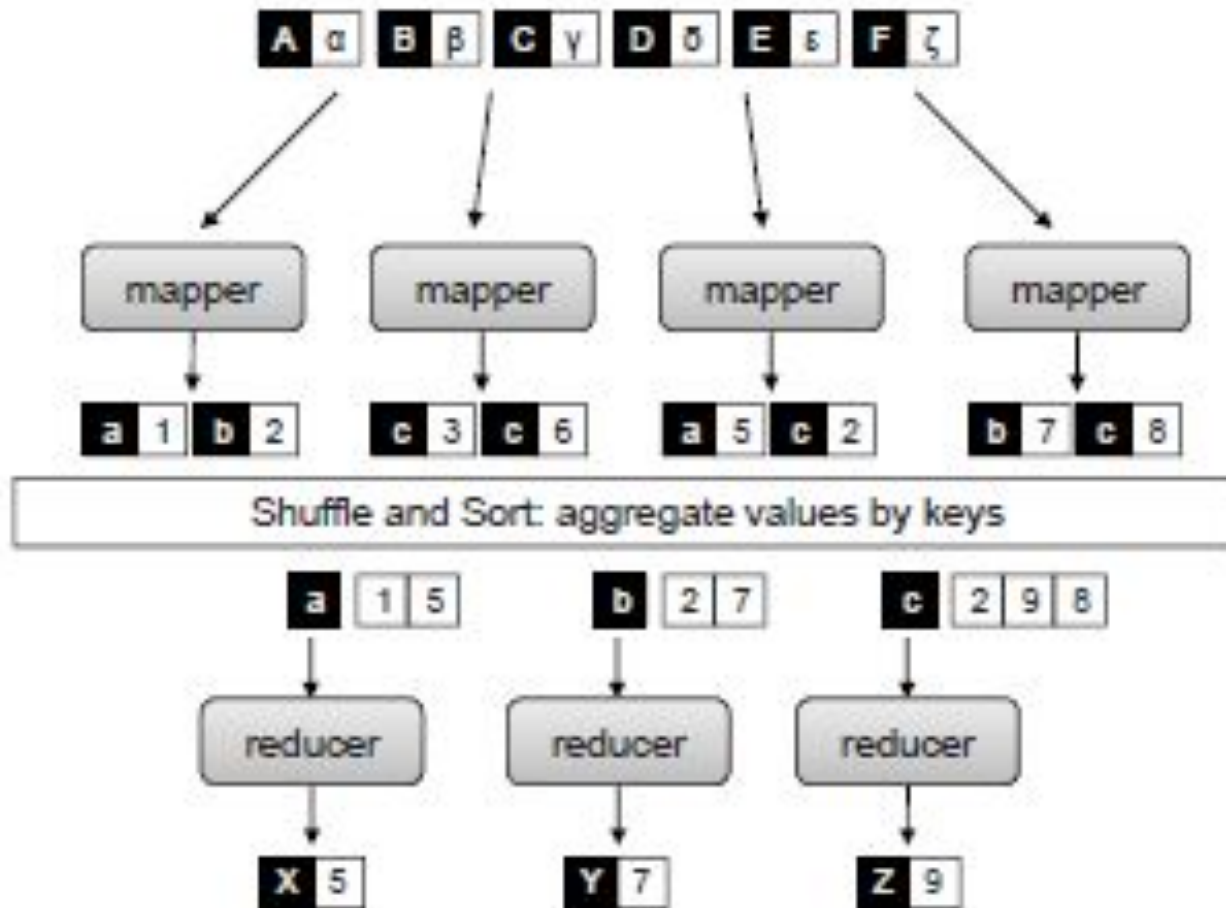
MapReduce Example

- Word count algorithm

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)
5:
6: class REDUCER
7:   method REDUCE(term  $t$ , counts [ $c_1, c_2, \dots$ ])
8:      $sum \leftarrow 0$ 
9:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
10:       $sum \leftarrow sum + c$ 
11:     EMIT(term  $t$ , count  $sum$ )
```

MapReduce Example

- Simplified MapReduce view



MapReduce

- Programmers specify two functions:
 - map** $(k, v) \rightarrow \langle k', v' \rangle^*$
 - reduce** $(k', v') \rightarrow \langle k', v' \rangle^*$
 - All values with the same key are sent to the same reducer
- The execution framework handles everything else...

What's “everything else”?

Execution Framework

- One of the most important ideas behind *MapReduce* is separating the ***what*** of distributed processing from the ***how***.
- A *MapReduce* program (job) consists of:
 - Code for *mappers* and *reducers* (as well as *combiners* and *partitioners* ... later)
 - Configuration parameters (e.g., where the input lies and where the output should be stored)
- The developer submits the job to the submission node of a cluster (*jobtracker*).
- The execution framework (*runtime*) takes care of everything else.

MapReduce *Runtime*

- **Handles scheduling**
 - Assigns workers to map and reduce tasks
- **Handles “data/code distribution”**
 - Moves processes to data
- **Handles synchronization**
 - Gathers, sorts, and shuffles intermediate data
- **Handles errors and faults**
 - Detects worker failures and restarts
- Everything happens on top of a distributed file system (GFS!)

Runtime Scheduling Issues

- ***Scheduling*** involves coordination among tasks belonging to different jobs.
- ***Speculative execution optimization*** (dealing with *Stragglers*)
 - The map phase of a job is only as fast as the slowest map task.
 - Completion of a job is bounded by the running time of the slowest reduce tasks.
 - Speculative execution makes identical copy of the same task and executes on a different machine, use first task completed.
 - Shown to improve running times by 44%.

Runtime Data/code Distribution Issues

- Need to somehow *feed data to co-located code*.
- Dependent on scheduling and underlying distributed file system.
- The scheduler starts tasks on the node that holds a particular block of data, i.e., local drive, needed by the task.
 - In effect moving the code to the data.
 - If this is not possible, it must stream data across the network.
 - An optimization here is to prefer nodes that are on the same rack in the datacenter – inter-rack BW is much slower.

Runtime Synchronization Issues

- **Synchronization** accomplished by a barrier between the *map* and *reduce* phases of processing.
 - Intermediate key-value pairs must be *grouped by key*.
 - Accomplished by *large distributed sort* involving all the nodes that executed map tasks and all the nodes that will execute reduce tasks.
 - Involves copying intermediate data over the network - “shuffle and sort”.
 - m mappers and r reducers involves up to $m \times r$ distinct copy operations.
 - Reduce computation can not start until all mappers have finished emitting key-value pairs & all key-value pairs have been shuffled and sorted!

How does distributed sort work?

Runtime Error and Fault Handling

- Disk failures are common
- RAM failures
- Planned and unplanned data center outages
- Software bugs
- Large data sets contain corrupted data
- *Rely heavily on distributed file system*

Partitioners and Combiners

- Two additional elements that complete the programming model: ***partitioners*** and ***combiners***.
- ***Partitioners*** – divide up intermediate key space and assign intermediate key-value pairs to reducers.
 - Within each reducer, keys are processed in sorted order
 - The simplest partitioner involves computing the hash value of the key and then taking the mod of that value with the number of reducers.

Partitioners and Combiners

- **Combiners** – optimization that allows for local aggregation before the shuffle and sort phase.

Word count example:

- Emits a key-value pair for each word in the collection!
- Key-value pairs need to be copied across the network!
- The amount of intermediate data will be larger than the input collection itself!!! – clearly not good.
- E.g., perform local aggregation on the output of each mapper, i.e., perform local count for a word over all the documents processed by the mapper.
- Number of words reduced to (at most) the number of unique words in the collection *times* the number of mappers – usually much smaller (Zipfian distribution of word occurrences).

MapReduce

- Programmers specify two functions:

map $(k, v) \rightarrow \langle k', v' \rangle^*$

reduce $(k', v') \rightarrow \langle k', v' \rangle^*$

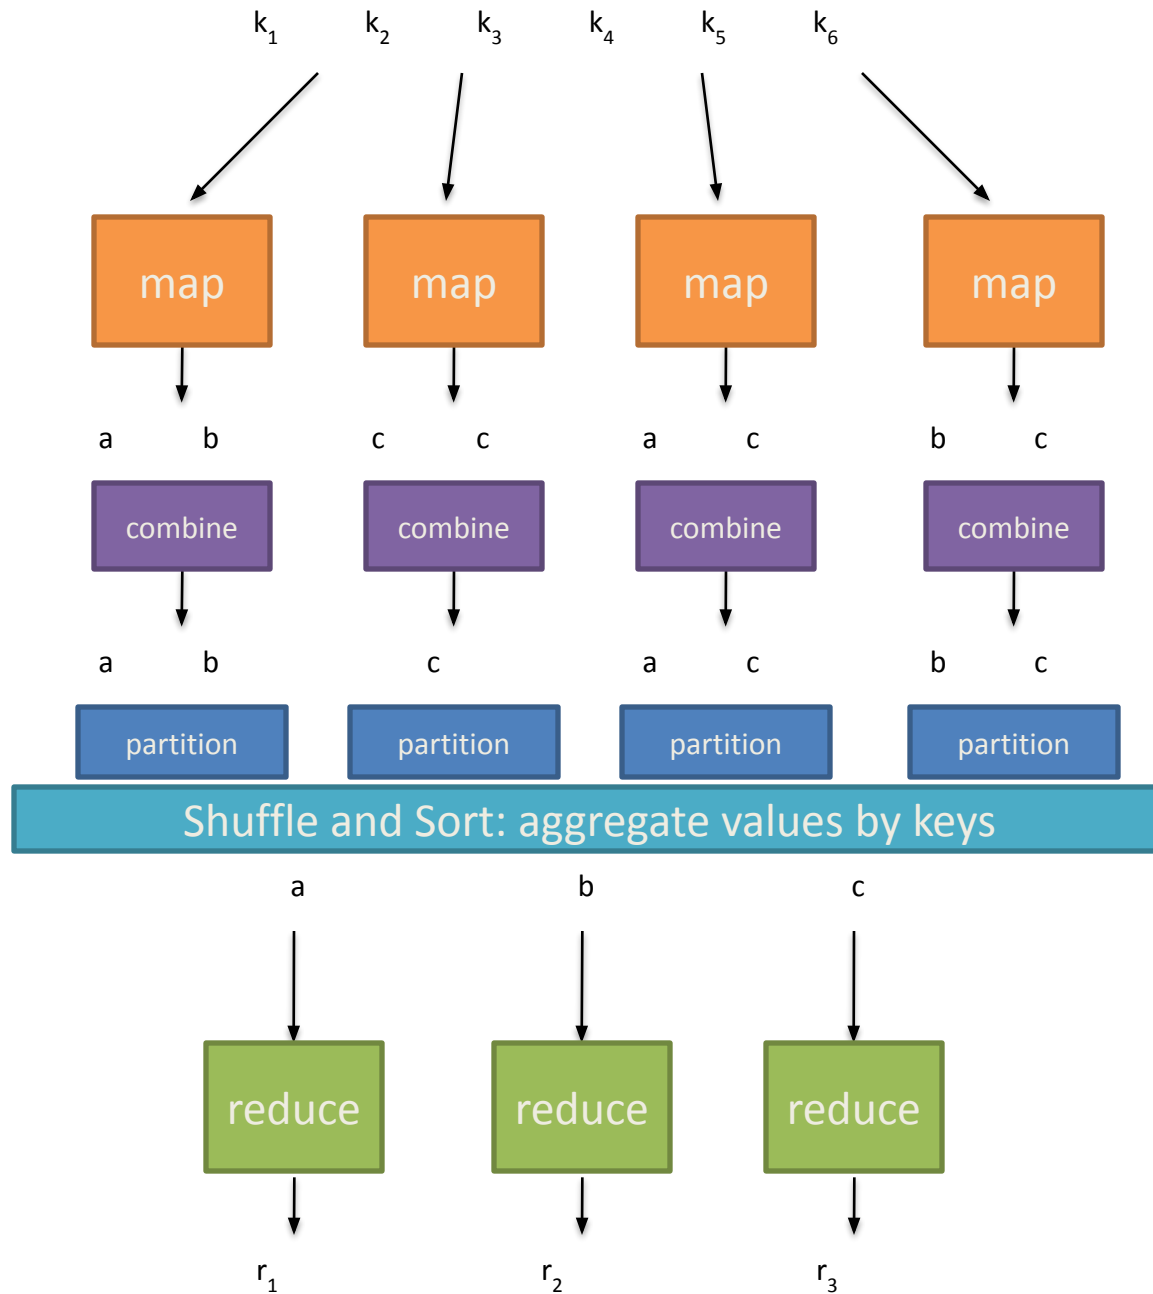
- All values with the same key are reduced together

combine $(k', v') \rightarrow \langle k', v' \rangle^*$

- Mini-reducers that run in memory after the map phase
- Used as an optimization to reduce network traffic

partition $(k', \text{number of partitions}) \rightarrow \text{partition for } k'$

- Often a simple hash of the key, e.g., $\text{hash}(k') \bmod n$
- Divides up key space for parallel reduce operations



How do we get data to the workers?



Distributed File System

- Don't move data to workers... move workers to the data!
 - Store data on the local disks of nodes in the cluster
 - Start up the workers on the node that has the data local
- Why?
 - Not enough RAM to hold all the data in memory
 - Disk access is slow, but disk throughput is reasonable
- A distributed file system is the answer
 - GFS (Google File System) for Google's MapReduce
 - HDFS (Hadoop Distributed File System) for Hadoop

GFS: Design Decisions

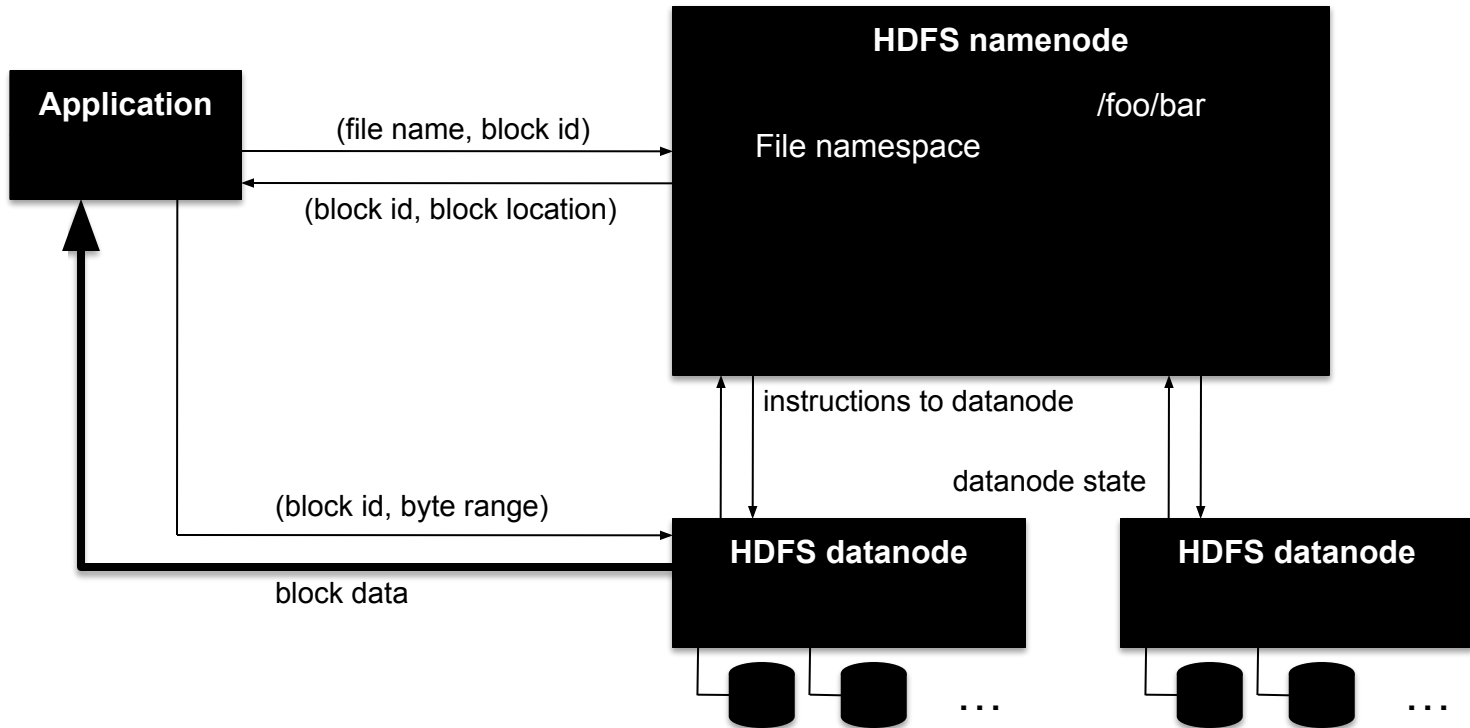
- Files stored as chunks
 - Fixed size (64MB)
- Reliability through replication
 - Each chunk replicated across 3+ chunk servers
- Single master to coordinate access, keep metadata
 - Simple centralized management
- No data caching
 - Little benefit due to large datasets, streaming reads
- Simplify the API
 - Push some of the issues onto the client (e.g., data layout)

HDFS = GFS clone (same basic ideas)

From GFS to HDFS

- Terminology differences:
 - GFS master = Hadoop namenode
 - GFS chunk servers = Hadoop datanodes
- Functional differences:
 - No file appends in HDFS
 - HDFS performance is (likely) slower

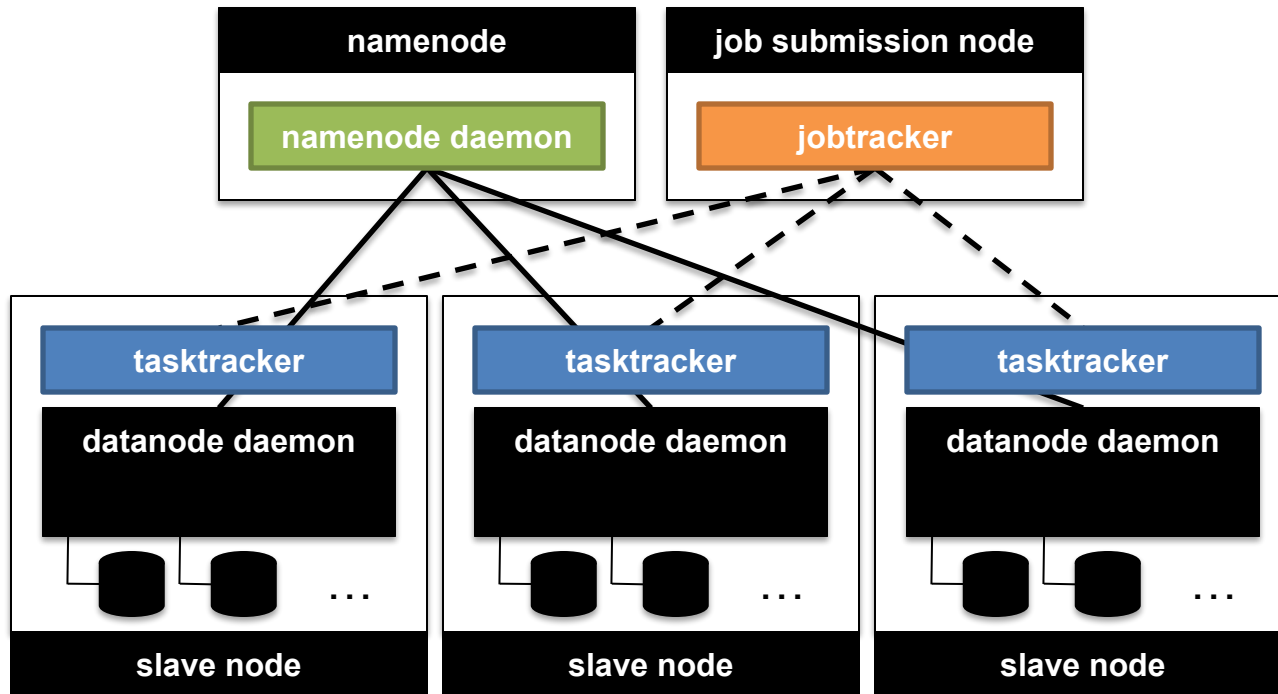
HDFS Architecture



Namenode Responsibilities

- Managing the file system namespace:
 - Holds file/directory structure, metadata, file-to-block mapping, access permissions, etc.
- Coordinating file operations:
 - Directs clients to datanodes for reads and writes
 - No data is moved through the namenode
- Maintaining overall health:
 - Periodic communication with the datanodes
 - Block re-replication and rebalancing
 - Garbage collection

Putting everything together...



Cloud Resources

- Hadoop on your local machine
- Hadoop in a virtual machine on your local machine
- Hadoop in the clouds with Amazon EC2
- Hadoop on the Google/IBM cluster

MapReduce Implementations

- Google has a proprietary implementation in C++
 - Bindings in Java, Python
- Hadoop is an open-source implementation in Java
 - Development led by Yahoo, used in production
 - Now an Apache project top level project
 - Rapidly expanding software ecosystem
- Lots of custom research implementations
 - For GPUs, cell processors, etc.